TBD

Wenrui Yuan, Stéphanie van den Berg, Johannes Steinrücke

March 9, 2026

Abstract

Knowledge tracing is a core component of adaptive learning systems, yet much of the literature assumes binary correctness as the observation signal. In many educational settings, however, learner evidence is polytomous, including ordered rubric levels and partial credit, where both ordinal structure and interpretability matter for measurement and intervention.

We propose DEEP-MIRT, a modular framework that extends knowledge tracing to polytomous responses by bridging sequential neural modeling with psychometric Item Response Theory. The framework couples a swappable sequential backbone with an interchangeable response head, separating the question of how learner ability evolves over time from the question of how that ability maps to ordered response categories. Instantiated with a Generalized Partial Credit Model (GPCM) head, DEEP-MIRT produces psychometrically interpretable outputs at every interaction, including an evolving student ability estimate, item discrimination values, and ordered difficulty thresholds. The model delivers strong predictive performance while accurately recovers ground-truth item parameters across varying category resolutions. The framework advances knowledge tracing from a purely predictive instrument into a measurement-aware system that grounds instructional decision-making in interpretable psychometric theory.

Introduction

Educational assessments often target graded performance rather than pass/fail outcomes. Rubric-based scoring, partial-credit items, and constructed-response tasks encode degrees of correctness that binary labels collapse into a single signal (Black and Wiliam, 1998; Heritage, 2007). While automated scoring systems can deliver this ordinal structure at scale in intelligent tutoring contexts (Shute, 2011; Baral et al., 2021), most KT benchmarks, evaluation protocols, and modeling paradigms still assume binary correctness. This leaves ordinal outcomes comparatively underexplored in KT, even though they are central to authentic assessment practice.

Beyond the outcome type, interpretability remains a key concern. Most KT models learn high-dimensional states optimized for prediction, but these states are not, by construction, tied to interpretable quantities. Therefore, constructs that educators and assessment designers routinely work with, such as student ability level, item discriminatory power, or the difficulty steps between adjacent rubric categories, are not directly

identifiable. Item Response Theory (IRT) provides exactly this structure through identifiable parameters including student ability, item discrimination, and ordered step difficulties, grounding score interpretation in established psychometric theory (Muraki, 1997; Embretson and Reise, 2000; Messick, 1995). Standard IRT models, however, treat ability as fixed and are not designed for sequential interaction data. This work therefore extends KT to polytomous responses by incorporating the interpretable parameter structure of IRT within a dynamic sequential framework.

Three lines of existing work address these issues in isolation, leaving a gap at their intersection. First, deep KT models, from Deep Knowledge Tracing (DKT; Piech et al., 2015) to memory-augmented and attention-based variants (Zhang et al., 2017; Pandey and Karypis, 2019), capture temporal learning dynamics but typically assume dichotomous correctness and learn latent states with limited psychometric interpretability. Second, polytomous Item Response Theory (IRT) models such as the Generalized Partial Credit Model (GPCM; Muraki, 1997) provide principled, interpretable measurement of ordered responses with identifiable parameters (student ability, item discrimination, ordered step thresholds), but they generally treat ability as time-invariant and are not designed for sequential interaction logs. Third, recent efforts that inject IRT structure into deep KT (Yeung, 2019; Converse et al., 2021) improve interpretability but largely remain restricted to dichotomous outcomes. An early partial exception is Wang and Heffernan (2013), who extended Bayesian Knowledge Tracing to allow partial credit via continuous Bayesian nodes, but without deep sequential modeling or recovering standard polytomous IRT step structure. Yet these approaches has not yet reconciled sequential dynamics with ordinal response modeling while retraining psychometric identifiability.

We propose *DEEP-MIRT*, a modular framework that extends KT to Polytomous Knowledge Tracing (PKT) by coupling a swappable sequential backbone with an IRT-grounded response head. In this setting, PKT is the task of predicting the full distribution over ordered response categories from a learner's interaction history. We instantiate the framework with a GPCM response head, yielding DEEP-GPCM, which combines a Dynamic Key-Value Memory Network (DKVMN) backbone with a GPCM observation model. At each timestep, the GPCM head maps the learner state to a distribution over ordinal categories and exposes psychometrically interpretable components, including an evolving student ability estimate, as well as learned item parameters such as per-item discrimination values and ordered step difficulties corresponding to rubric level transitions. While the experiments focus on ordinal PKT, the framework accommodates other polytomous formats through alternative response heads.

Our main contributions are:

1. **A modular framework for polytomous knowledge tracing.** We formulate PKT as predicting the full ordinal response distribution from interaction history and propose DEEP-MIRT, which decouples a sequential backbone (encoder, memory, summary) from an interchangeable psychometric response head. The head interface accepts student ability θ_t and item identity $\mathbf{k}_t$ and returns a category probability vector; any differentiable observation model satisfying this contract is a valid instantiation. This design separates the educational measurement question (which model governs how ability maps to rubric scores?) from the knowl-

edge tracing question (how does ability evolve over time?).

2. **GPCM instantiation with interpretable psychometric parameters.** We instantiate DEEP-MIRT with a GPCM-based response head as DEEP-GPCM, incorporating a monotonic threshold parameterization that enforces rubric ordering by construction. The resulting model produces evolving ability trajectories, item discrimination estimates, and ordered step difficulty thresholds, all of which correspond to constructs that educators and assessment designers already use in rubric design and standard-setting practice.

3. **Scalable item encoding and empirical evaluation.** We propose Static Item Embedding (SIE) for large item banks and evaluate ordinal prediction quality, psychometric parameter recovery, ablation of design components, and backward compatibility with binary KT benchmarks.

We investigate the following research questions:

- **RQ1 (Prediction).** Does coupling a DKVMN backbone with an ordinal, psychometrically constrained response head improve polytomous next-response prediction, and does it better preserve ordinal structure than format-unaware alternatives?

- **RQ2 (Learner State).** Do the inferred student ability trajectories θ_t exhibit psychologically plausible learning dynamics, and can DEEP-GPCM reveal meaningful learner patterns (mastery, growth, persistent difficulty) that are not visible from format-unaware latent state representations?

- **RQ3 (Parameter Recovery).** Do the IRT parameters inferred by DEEP-GPCM faithfully recover known ground-truth values from synthetic data, and are item parameter estimates stable across student subpopulations as assessed by split-half reliability?

- **RQ4 (Scalability).** Does Static Item Embedding maintain ordinal prediction quality and parameter recovery at large item-bank scales ($Q \sim 5000$) while substantially reducing parameter count compared to dense encoding strategies?

- **RQ5 (Class Imbalance).** Does DEEP-GPCM maintain ordinal prediction quality and IRT parameter recovery under realistic class imbalance, where higher rubric scores dominate the response distribution as is common in operational educational assessments?

Related Work

Knowledge Tracing: From Binary to Polytomous

Knowledge tracing has evolved through several modeling paradigms, but most of them assume a binary mastery state. Bayesian Knowledge Tracing (BKT) models mastery of a Knowledge Component (KC) as a two-state Hidden Markov Model (Corbett and

Anderson, 1994). Deep Knowledge Tracing (DKT) reframed the task as sequence modeling with recurrent neural networks, learning high-dimensional latent states from interaction sequences (Piech et al., 2015). Dynamic Key-Value Memory Networks (DKVMN) introduced a structured memory that separates static concept representations (key memory) from evolving mastery states (value memory), offering a more transparent mapping between knowledge components and latent state than generic RNNs (Graves et al., 2016; Zhang et al., 2017). Attention-based models such as SAKT and AKT apply self-attention to weigh the relevance of prior interactions for next-step prediction (Pandey and Karypis, 2019; Ghosh et al., 2020), and recent work has further explored Transformer-based architectures (Choi et al., 2020; Liu et al., 2023). Throughout this progression, the response format has remained dichotomous: every model cited above predicts a binary correct/incorrect outcome.

Recent KT research has begun to acknowledge that correctness alone is insufficient. Lu et al. (2024) fuse process features (attempts, response time, historical accuracy) to enrich difficulty representations, and Li et al. (2025) integrate problem complexity and state-stability mechanisms to reduce noisy updates. These efforts treat richer signals as auxiliary features within a binary prediction framework. They do not change the response format itself, but they highlight that educational interactions carry information beyond pass/fail: precisely the information that ordinal and polytomous responses encode directly. The only prior work to move beyond binary responses in a KT context is Wang and Heffernan (2013), who extended BKT to partial credit using continuous Bayesian nodes, but without deep sequential modeling or IRT-based parameterization.

Interpretability in Deep Knowledge Tracing

In most DKT models, learned states are computed as high dimensional representations than psychometrically meaningful constructs (National Research Council, 2001). Several efforts address this gap by coupling neural architectures with Item Response Theory (IRT) structure. Deep-IRT attaches an IRT-style readout to DKVMN, exposing student ability and item difficulty parameters for each prediction (Yeung, 2019). DKVMN-MRI similarly integrates IRT within an LSTM-based memory framework (Xu et al., 2024). Converse et al. (2021) propose a general approach for incorporating IRT parameters into KT training. These IRT-augmented models improve interpretability: a teacher or assessment specialist can inspect the estimated ability and difficulty values, rather than relying on opaque hidden states. In educational terms, "interpretability" here means that model outputs correspond to constructs that practitioners already use (student proficiency levels, item difficulty, and discrimination quality), supporting the kinds of validity arguments that responsible assessment requires (Messick, 1995; Kane, 2013).

A critical limitation, however, is that all existing IRT-KT integrations operate on dichotomous responses. They estimate a single ability and difficulty for binary outcomes and do not model the ordered step structure of rubric scores or partial-credit items. Attention-based models provide a complementary form of interpretability through attention weights (which interactions were most relevant to a prediction), but they too lack psychometric parameterization and ordinal structure awareness.

Polytomous Item-Response Theory

Psychometrics provides a family of models for ordered responses grounded in a latent ability parameter θ and item-specific properties (Hambleton and Swaminathan, 1985; Embretson and Reise, 2000). Two traditions dominate for ordinal outcomes. Cumulative models such as the Graded Response Model (GRM; Samejima, 1969) define boundary parameters between adjacent categories, modeling the probability of scoring at or above category k as a logistic function of θ. Adjacent-category models, including the Partial Credit Model (PCM; Masters and Wright, 1997) and the Generalized Partial Credit Model (GPCM; Muraki, 1997), parameterize each step transition directly: the GPCM characterizes each item by a discrimination parameter α_j controlling how sharply the item differentiates ability levels, and ordered step difficulty thresholds $\boldsymbol{\beta}_j = (\beta_{j,0}, \ldots, \beta_{j,K-2})$ representing the ability level at which each consecutive rubric transition becomes equally likely. Adjacent-category models are the natural fit for rubric scoring, where performance progresses sequentially through ordered levels (Embretson and Reise, 2000; Wilson, 2005). For unordered categories, the Nominal Response Model (NRM; Bock, 1972) provides an analogous framework.

In the context of knowledge tracing, these static measurement paradigms must be reformulated to capture sequential learning dynamics. Formally, at time step t, a student interacts with an item indicated by a one-hot vector $q_t \in \{0,1\}^Q$ over Q unique items, yielding an ordinal response $r_t \in \{0, 1, \ldots, K-1\}$ across K rubric levels. Defining an interaction as the tuple $x_t = (q_t, r_t)$ and the cumulative history as $\mathcal{X}_t = \{x_1, \ldots, x_t\}$, the objective of polytomous knowledge tracing is to predict the conditional distribution of the next response:

$$P(r_{t+1} = k \mid q_{t+1}, \mathcal{X}_t), \quad k \in \{0, \ldots, K-1\}. \tag{1}$$

This work therefore explores how the interpretable parameter structures of polytomous IRT might be adapted for sequential modeling. We treat it as a time-varying latent state conditioned on the interaction history $\mathcal{X}_t$, aiming to interpret dynamic tracing results with psychometric models.

Methodology

Framework Overview

DEEP-MIRT is a modular architecture with two separable concerns. The *backbone* handles the knowledge tracing problem: it encodes item–response interactions, maintains a dynamic memory of the student's evolving knowledge state, and produces a summary vector $\mathbf{f}_t$ representing current student state. The *response head* handles the measurement problem: it maps the backbone's student state $\mathbf{f}_t$ and item identity $\mathbf{k}_t$ to a probability distribution over response categories through a psychometric observation model. Any differentiable function $h(\mathbf{f}_t, \mathbf{k}_t) \rightarrow \Delta^{K-1}$ (a distribution over K categories) satisfying this interface is a valid DEEP-MIRT instantiation. The backbone is agnostic to which measurement model governs the head.

This separation is conceptually and practically helpful for tracking learning progressions. The backbone addresses the learning dynamics question of how student

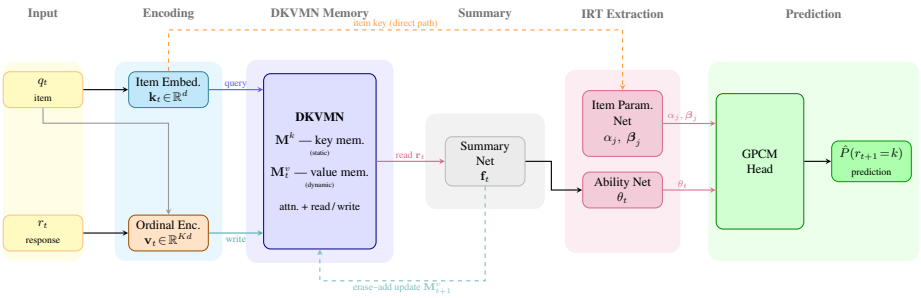

Figure 1: DEEP-MIRT framework at timestep t. Item parameters $(\alpha_j, \boldsymbol{\beta}_j)$ bypass the student state via the item embedding (dashed, top), consistent with the IRT assumption that discrimination and difficulty are student-independent. The GPCM head maps $(\theta_t, \alpha_j, \boldsymbol{\beta}_j)$ to the predicted ordinal distribution $\hat{P}(r_{t+1} = k)$.

ability evolves across a sequence of interactions, naturally handled by sequential neural architectures. The head addresses the measurement question of how current ability maps to the probability of each ordered rubric score, naturally handled by psychometric theory. Coupling them through a clean interface means the two concerns can be improved independently: a better backbone (e.g., Transformer instead of DKVMN) improves temporal modeling without changing the interpretable parameterization; a different head (e.g., GRM instead of GPCM) changes the measurement assumptions without retraining the backbone.

In this paper we instantiate the head with the GPCM over the other major polytomous families for three reasons: (1) it is an adjacent-category model, which is the natural match for rubric scoring where a student progresses sequentially through performance levels; (2) it produces a discrimination parameter α_j and ordered step difficulty thresholds $\boldsymbol{\beta}_j$ that correspond directly to rubric design decisions and can be inspected by assessment practitioners; (3) its monotone threshold constraint admits a clean differentiable parameterization as described below. A cumulative model such as the GRM (Samejima, 1969) would be a natural alternative for tasks where each scoring threshold represents an independent decision (e.g., analytic rubrics with independent sub-scores); a NRM (Bock, 1972) would apply when response categories are unordered (e.g., multiple-choice representing different misconceptions). We treat these as directions for future instantiation.

The backbone processes inputs through four sequential stages as shown in Figure 1: ordinal interaction encoding, DKVMN memory update, summary compression, and IRT-based parameter extraction. An ordinal-aware training objective supervises all stages jointly. The following subsections formalize each component.

Ordinal-Aware Interaction Representation

Linear Decay Encoding We encode an ordinal response r_t with a "triangular" kernel over ordered categories. Define a length-K ordinal code $\mathbf{s}_t \in \mathbb{R}^K$ by

$$\begin{aligned} s_t^{(k)} &= \max\left(0, 1 - \frac{|k - r_t|}{K-1}\right), \quad 0 \le k \le K-1, \\ \mathbf{s}_t &= [s_t^{(0)}; \ldots; s_t^{(K-1)}]. \end{aligned} \tag{2}$$

For example, with a $K = 4$ rubric and a response of $r_t = 2$ ("correct approach with errors"), the ordinal code is $\mathbf{s}_t = [0.33,\ 0.67,\ 1.0,\ 0.67]$, which spreads activation to neighboring categories while peaking at the observed level. This encoding preserves the ordinal structure: a response of 2 is represented as more similar to responses 1 and 3 than to response 0, reflecting the educational intuition that partial credit is closer to full credit than to no credit.

We then bind the ordinal code to the item indicator via a Kronecker product (equivalently, a block one-hot expansion):

$$\mathbf{x}_t = q_t \otimes \mathbf{s}_t \in \mathbb{R}^{KQ}. \tag{3}$$

This reduces to the standard binary interaction encoding when $K = 2$ (consistent with Zhang et al. (2017)). While this encoding is effective for moderately sized item banks, the $\mathcal{O}(KQ)$ dimensionality becomes a practical constraint for large Q; we address this in the following subsection.

Scalable Value Encoding for Large Item Banks Large-scale educational platforms maintain extensively large item banks with various grade levels. The triangular Kronecker encoding in (2) yields an interaction vector of dimension KQ, which grows linearly with item bank size. The subsequent value projection $\mathbf{v}_t = B^\top \mathbf{x}_t \in \mathbb{R}^{d_v}$ therefore requires a $KQ \times d_v$ weight matrix; at $Q = 5000$ and $K = 5$, this amounts to 1.6M parameters and a per-batch intermediate tensor of size $\mathrm{batch} \times S \times 25000$, placing substantial memory pressure on GPU-constrained deployment environments.

We propose **Static Item Embedding** (SIE), an alternative value encoding that decouples item identity from response pattern via two independent additive branches:

$$\mathbf{v}_t = W_{\text{item}}\, \mathbf{u}_{q_t} + W_{\text{resp}}\, \mathbf{s}_t + \mathbf{b}_{\text{resp}} \in \mathbb{R}^{d_v}, \tag{4}$$

where $\mathbf{u}_{q_t} \in \mathbb{R}^H$ is a *frozen* (non-trainable) unit-norm random vector assigned to item q_t at initialization, $\mathbf{s}_t \in \mathbb{R}^K$ is the triangular ordinal code from (2), and $W_{\text{item}} \in \mathbb{R}^{d_v \times H}$ and $W_{\text{resp}} \in \mathbb{R}^{d_v \times K}$ are learned projections shared across all items.

The factored structure in (4) separates two gradient channels: W_{item} is updated by item-identity signals, while W_{resp} is updated by response-pattern signals. Because item vectors are frozen, items cannot collapse to similar representations during training, a failure mode of learned separable embeddings at large Q (as shown in the ablation study below). For two independent uniformly random unit vectors in $\mathbb{R}^H$, the expected squared inner product is $1/H$ (Dasgupta and Gupta, 2003), so pairwise interference

concentrates around zero as H grows; cross-item gradient crosstalk in W_{item} therefore decays with H independently of Q, making SIE well-suited for large item banks.

The frozen dimension H is set as the smallest power of two no less than $\max(128, \lfloor Q/2 \rfloor)$, clamped to $[128, 1024]$.[1] This schedule ensures near-orthogonality scales with item bank size while keeping the projection dimension bounded, in contrast to the $\mathcal{O}(KQ)$ complexity of the full Kronecker encoding.

Dynamic Key-Value Memory Network

We adopt the Dynamic Key-Value Memory Network (DKVMN) as the sequential backbone. Let $\mathbf{k}_t = A^\top q_t \in \mathbb{R}^{d_k}$ be a key embedding for the current item and $\mathbf{v}_t \in \mathbb{R}^{d_v}$ be a value embedding for the interaction, obtained via the Kronecker projection $\mathbf{v}_t = B^\top \mathbf{x}_t$ (ordinal-aware encoding) or the Static Item Embedding in (4) (scalable encoding). The model maintains a static key memory $\mathbf{M}^k \in \mathbb{R}^{N \times d_k}$ and a dynamic value memory $\mathbf{M}_t^v \in \mathbb{R}^{N \times d_v}$. Educationally, the key memory represents latent skill areas or knowledge components, the value memory tracks the student's evolving mastery of each area, and the attention weights $\mathbf{w}_t$ (below) indicate which skill areas a given item exercises. The summary vector $\mathbf{f}_t$ aggregates the student's current mastery state relevant to the item at hand.

Specifically, the key query is first projected via a learned transform, and attention weights over memory slots are computed as

$$\tilde{\mathbf{k}}_t = \tanh(W_{\text{attn}} \mathbf{k}_t), \qquad \mathbf{w}_t = \text{softmax}(\mathbf{M}^k \tilde{\mathbf{k}}_t) \in \mathbb{R}^N, \qquad w_t(i) = [\mathbf{w}_t]_i, \quad (5)$$

and the read vector is

$$\mathbf{r}_t = \sum_{i=1}^{N} w_t(i)\, \mathbf{M}_t^v(i) \in \mathbb{R}^{d_v}. \quad (6)$$

We fuse item context and retrieved state into a summary vector

$$\mathbf{f}_t = \tanh(W_{\mathbf{f}} [\mathbf{r}_t; \mathbf{k}_t] + b_{\mathbf{f}}), \quad (7)$$

which is then used by the response head to predict r_{t+1}.

Following memory-augmented neural networks, the value memory is updated via erase-add operations:

$$\mathbf{e}_t = \sigma(W_{\mathbf{e}} \mathbf{v}_t + b_{\mathbf{e}}), \quad (8)$$

$$\mathbf{a}_t = \tanh(W_{\mathbf{a}} \mathbf{v}_t + b_{\mathbf{a}}), \quad (9)$$

$$\mathbf{M}_{t+1}^v(i) = \mathbf{M}_t^v(i)\big[1 - w_t(i)\mathbf{e}_t\big] + w_t(i)\mathbf{a}_t. \quad (10)$$

DEEP-GPCM: an Ordinal IRT Head

Instead of the binary readout (e.g., $\sigma(W_p^\top \mathbf{f}_t + b_p)$), we use a Generalized Partial Credit Model (GPCM) head to obtain category probabilities for ordered responses. Let j denote the item exercised at time t (i.e., $[q_t]_j = 1$). We compute a time-varying ability

[1] Concretely, $H = 128$ for $Q \leq 256$, $H = 256$ for $Q \approx 500$, $H = 512$ for $Q \approx 1000$, and $H = 1024$ for $Q \geq 2000$.

proxy θ_t and item-specific parameters (thresholds and discrimination) as

$$\theta_t = W_\theta \mathbf{f}_t + b_\theta, \tag{11}$$

$$\beta_{j,0} = W_\beta \mathbf{k}_t + b_\beta, \tag{12}$$

$$\beta_{j,k} = \beta_{j,k-1} + \mathrm{softplus}(W_{g[k]}\mathbf{k}_t + b_{g[k]}), \quad 1 \le k \le K-2, \tag{13}$$

$$\alpha_j = \exp(W_\alpha \mathbf{k}_t + b_\alpha)\,. \tag{14}$$

The construction in (13) enforces $\beta_{j,0} < \beta_{j,1} < \cdots < \beta_{j,K-2}$ by adding strictly positive increments (via softplus) to each preceding threshold. This monotonicity guarantee has direct educational meaning: achieving rubric level $k{+}1$ always requires higher ability than achieving level k, respecting the ordering that a rubric designer intended. Without this constraint, the model could learn disordered thresholds where some rubric levels are never the most probable response at any ability level, a well-documented problem in educational measurement that signals a poorly functioning item (Wilson, 2005).

Because discrimination is an item-intrinsic property in the GPCM, α_j is derived solely from the item key $\mathbf{k}_t$, deliberately excluding the student-state summary $\mathbf{f}_t$. This separation ensures that the estimated discrimination reflects the item's inherent capacity to differentiate ability levels, independently of which student is currently being modelled — consistent with the IRT assumption that item parameters are population-invariant (Hambleton and Swaminathan, 1985). Item-level discrimination is recovered by averaging per-timestep estimates across all encounters with a given item (Meulders and Xie, 2004). The exponential activation ensures $\alpha_j > 0$.

Given $(\theta_t, \alpha_j, \{\beta_{j,k}\}_{k=0}^{K-2})$, the GPCM category probability is

$$p_{t,k} = \frac{\exp\Big(\sum_{i=0}^{k-1} \alpha_j(\theta_t - \beta_{j,i})\Big)}{\sum_{c=0}^{K-1} \exp\Big(\sum_{i=0}^{c-1} \alpha_j(\theta_t - \beta_{j,i})\Big)}. \tag{15}$$

At each timestep, the model evaluates this standard GPCM formulation using the *dynamically estimated* student ability θ_t and item parameters $\alpha_j, \boldsymbol{\beta}_j$. Unlike classical GPCM estimation where ability is static, here θ_t evolves as new interactions are observed, enabling the model to track learning over time while retaining the interpretable measurement structure of polytomous IRT.

We predict $\hat{r}_{t+1} = \arg\max_k p_{t,k}$, and optimize the weighted ordinal loss in (16) over the observed responses.

Training Objective for Ordinal Data

Because the GPCM defines category probabilities via (15) rather than cumulative links, we optimize a weighted ordinal loss (WOL) that jointly addresses class imbalance and ordinal distance (Gutiérrez et al., 2016):

$$\mathcal{L}_{\mathrm{WOL}} = \frac{1}{n}\sum_{i=1}^{n}\Big[(1 + \gamma|y_i - \hat{y}_i|) \cdot w_{y_i} \cdot (-\log \hat{p}_{i,y_i})\Big], \tag{16}$$

where y_i is the true category, $\hat{p}_{i,y_i}$ is the predicted probability under (15), $\hat{y}_i = \arg\max_k \hat{p}_{i,k}$ is the predicted label, and γ controls the ordinal distance penalty. To balance classes, with c_k the sample count of class k, we set $w_k = \sqrt{\frac{n}{K \cdot c_k}}$, where n is the number of valid tokens in the batch. The distance term $(1 + \gamma|y_i - \hat{y}_i|)$ penalizes predictions that are far from the true category on the ordinal scale, encouraging the model to prefer near-misses over large errors—a property that standard cross-entropy does not enforce.

Experiments

Datasets

Synthetic-Ordinal. To the best of our knowledge, there are no standardized public benchmarks for ordinal knowledge tracing. We therefore generate a synthetic dataset using an IRT-GPCM data generating process following established psychometric principles. This design provides controlled polytomous supervision $r_t \in \{0, \ldots, K-1\}$ and enables parameter recovery analysis with known ground truth.

We simulate 5000 students interacting with questions drawn from a pool of $Q = 200$ questions in random order. Each interaction sequence has variable length in $[20, 80]$. Students and questions are configured with *static* GPCM parameters following assumed priors:

- *Student ability:* $\theta \sim \mathcal{N}(0, 1)$,
- *Item discrimination:* $\alpha \sim \mathrm{LogNormal}(0, 0.3)$,
- *Item thresholds:* $K-1$ values drawn independently from $\mathcal{N}(0, 1)$ and sorted to obtain $\beta_0 < \cdots < \beta_{K-2}$.

We generate four datasets with $K \in \{2, 3, 4, 5\}$ to test extensibility across category resolutions. Each dataset is split 80/20 for training/testing; a validation split is held out from training for early stopping and hyperparameter tuning.

Synthetic-Imbalanced. Real educational assessments rarely produce uniformly distributed rubric scores. High-performing students dominate most operational datasets, creating skewed response distributions where the highest category is most frequent. To test robustness to this common condition, we generate three imbalanced variants of the $Q = 200$, $K = 4$ dataset by shifting the student ability distribution upward: mild imbalance ($\theta \sim \mathcal{N}(0.5, 1)$), severe imbalance ($\theta \sim \mathcal{N}(1.0, 1)$), and extreme imbalance ($\theta \sim \mathcal{N}(1.5, 1)$). This produces progressively skewed response distributions where the highest category accounts for approximately 35%, 50%, and 65% of responses respectively, while retaining the same item parameters and GPCM data-generating process. All other settings match Synthetic-Ordinal.

Binary Compatibility After validating the PKT instantiations on Synthetic-Ordinal, we confirm compatibility with binary KT. Due to implementation differences in batch handling and sequence-to-question lookup, we select two datasets:

- **ASSIST2015**[2]: 19,840 students and 100 KCs/questions.
- **Synthetic-5**[3]: (Piech et al., 2015) simulated 4,000 students' responses with probability $p(r) = c + \frac{1}{1+\exp(-(\theta-\beta))}$ where $c = 0.25$ models question guessing. Each student is administered 50 questions in sequential order.

Models and Baselines

We compare DEEP-GPCM against three baselines that progressively relax psychometric and ordinal structure, each isolating a different design question:

- **DKVMN + Softmax**: DKVMN backbone with a K-way softmax head trained with categorical cross-entropy. No ordinal structure, no IRT constraints. Tests whether psychometric parameterization provides gains over a purely neural polytomous baseline.
- **Static GPCM**: Standard GPCM fit from interaction data with student ability held fixed across items. No sequential dynamics. Represents the psychometric status quo and quantifies what sequential modeling adds beyond classical IRT estimation.
- **Dynamic GPCM**: GPCM with per-item α/β lookup tables and a gated recurrent ability update driven by prediction error at each step. No DKVMN memory. Isolates the contribution of sequential ability updating from DKVMN's distributed memory.
- **DEEP-GPCM (ours)**: DKVMN backbone with GPCM head and weighted ordinal loss (16). Uses SIE encoding by default; the choice of encoding strategy is ablated below.

Prediction Mapping

For $K = 2$, predicted probabilities are converted to labels with a threshold of 0.5 following Zhang et al. (2017). For $K > 2$, we use $\hat{y}_t = \arg\max_k \hat{p}_{t,k}$ to map categorical probabilities to prediction labels.

Training Setup

We use the weighted ordinal loss in (16) to handle class imbalance and ordinal agreement.

We set $d_k = 64$ and $d_v = 128$ for key/value embedding dimensions and configure $N = 50$ memory slots for both memory matrices.

All models are trained with batch size 64 using Adam with learning rate 10^{-3}, reduced on performance plateau as per Mukherjee et al. (2019) for late-epoch stability. We apply dropout 0.05 and early stopping if validation loss does not improve for 10 epochs. Models are implemented in PyTorch and trained on an RTX 4060 Laptop GPU.

[2]https://sites.google.com/site/assistmentsdata/datasets/2015-assistments-skill-builder-data
[3]https://github.com/chrispiech/DeepKnowledgeTracing/tree/master/data/synthetic

Evaluation Metrics

We use the following metrics:

- **Categorical Accuracy (ACC)**: percentage of exact matches between predicted category and ground truth. While straightforward, ACC does not distinguish near-misses from far-misses and is therefore insufficient as a sole ordinal metric.
- **Quadratic Weighted Kappa (QWK)** (Cohen, 1968): our primary metric for ordinal tasks. QWK applies a quadratic penalty proportional to the squared distance between predicted and true categories, making it the standard agreement measure in automated essay scoring and rubric evaluation. A QWK of 0.0 indicates chance-level agreement; 1.0 indicates perfect agreement.
- **Kendall's Tau (τ)**: measures rank-order agreement between predicted and true response sequences, penalizing all inversions equally regardless of distance.
- **Area Under the ROC Curve (AUC)**: used for the binary case ($K = 2$) to maintain comparability with prior KT literature.

IRT Analysis Setup

For parameter recovery analysis on Synthetic-Ordinal, we use:

- $\boldsymbol{\theta}$ **Aggregation**: we take the mean ability estimate $\bar{\theta} = \frac{1}{T}\sum_{t=1}^{T}\theta_t$ across all timesteps for each student, averaging out sequence-position noise while exploiting the full interaction history.
- $\boldsymbol{\alpha}, \boldsymbol{\beta}$ **Aggregation**: because questions appear with variable frequency due to random sequencing, we aggregate item parameters based on occurrences and use the sample mean for recovery analysis.

To enable fair comparison between estimated and true parameters, we apply IRT linking transformations that align the scales while preserving correlations. For student ability θ and item difficulty β, we use z-score normalization: $\theta_{\text{norm}} = (\theta - \mu_\theta)/\sigma_\theta$, where μ and σ are computed from the respective parameter sets (estimated or true). For discrimination α, we apply log-transformation followed by standardization: $\alpha_{\text{norm}} = \exp\big((\log\alpha - \mu_{\log\alpha})/\sigma_{\log\alpha} \cdot 0.3\big)$, where the 0.3 scaling factor matches the log-normal prior used in data generation ($\alpha \sim \text{LogNormal}(0, 0.3)$). This normalization is standard practice in IRT parameter recovery studies, as it removes arbitrary scale and location differences while preserving the correlation structure that reflects true parameter recovery quality.

Results

Ordinal Prediction on Synthetic-Ordinal (RQ1)

Table 1 shows that DEEP-GPCM consistently outperforms all baselines across category resolutions. DEEP-GPCM achieves the highest QWK at every K, with the advantage

becoming more pronounced as the number of categories increases: at $K = 3$, DEEP-GPCM achieves QWK of 0.592 versus 0.578 for DKVMN+Softmax; at $K = 6$, the gap widens to 0.791 versus 0.773. This pattern demonstrates that the GPCM head's explicit ordinal structure becomes increasingly valuable as rubric complexity grows. DEEP-GPCM also outperforms both psychometric baselines: Dynamic GPCM achieves competitive QWK (0.704 at $K = 6$) but falls short of the neural models, while Static GPCM's lack of sequential dynamics results in poor prediction quality (QWK < 0.18 across all K).

Table 1: Ordinal prediction on Synthetic-Ordinal ($Q = 200$). DEEP-GPCM outperforms all baselines. Best QWK per K in bold.

	$K = 3$		$K = 4$		$K = 5$		$K = 6$	
Model	**ACC**	**QWK**	**ACC**	**QWK**	**ACC**	**QWK**	**ACC**	**QWK**
Static GPCM	.407	.168	.310	.156	.262	.178	.216	.150
Dynamic GPCM	.549	.533	.459	.619	.389	.681	.337	.704
DKVMN+Softmax	.597	.578	.531	.667	.511	.740	.476	.773
DEEP-GPCM	**.596**	**.592**	**.529**	**.682**	**.507**	**.759**	**.470**	**.791**

The key finding is that coupling a DKVMN backbone with a psychometrically constrained response head improves ordinal prediction quality compared to format-agnostic alternatives. DKVMN+Softmax treats ordinal categories as arbitrary labels, while DEEP-GPCM's GPCM head enforces the assumption that categories represent ordered performance levels. This structural prior guides learning toward solutions that respect rubric ordering, yielding better generalization on ordinal prediction tasks. The results demonstrate that interpretability and prediction quality are not in conflict: psychometric structure can improve both simultaneously.

Binary Compatibility on Standard KT Benchmarks

The GPCM head collapses to logistic regression for binary responses, enabling backward compatibility. Table 2 shows that DEEP-GPCM achieves competitive performance with DKT and DKVMN on real benchmarks (ASSIST2015, Synthetic-5). Psychometric parameterization does not sacrifice predictive quality.

Table 2: Binary compatibility ($K = 2$). DEEP-GPCM matches neural baselines. Best per column in bold.

Model	**Synthetic-Ordinal**		**ASSIST2015**		**Synthetic-5**	
	ACC	**QWK**	**ACC**	**AUC**	**ACC**	**AUC**
DKT	—	—	75.2%	0.727	**75.2**%	0.817
DKVMN	—	—	75.1%	**0.729**	75.1%	**0.827**
DEEP-GPCM	**70.6**%	**0.410**	**75.4**%	0.723	74.7%	0.824

Learner State Dynamics (RQ2)

Figure 2 visualizes estimated ability trajectories θ_t for four learner archetypes: high-ability learners who consistently perform well, low-ability learners who struggle throughout, improving learners who show growth over time, and declining learners whose performance deteriorates. The figure reveals three key findings. First, DEEP-GPCM and DKVMN+Softmax produce nearly identical trajectories across all archetypes, confirming that the GPCM head preserves the memory network's learned dynamics while reinterpreting them in psychometric terms. The GPCM head acts as a measurement lens that exposes interpretable ability estimates θ_t and item parameters (α, β) without altering the underlying sequential learning signal. Second, DEEP-GPCM successfully differentiates learner archetypes: high-ability trajectories remain elevated ($\theta \approx 1.5$), low-ability trajectories remain depressed ($\theta \approx -1.5$), and growth/decline patterns show clear temporal trends. This differentiation is essential for adaptive testing applications, where item selection must respond to real-time ability estimates. Third, Dynamic GPCM's trajectories remain flat regardless of archetype, hovering near $\theta \approx 0$ for all learners. This failure reveals that recurrent aggregation alone cannot disentangle individual ability from item characteristics—distributed memory is necessary to maintain learner-specific state representations that evolve meaningfully over time.

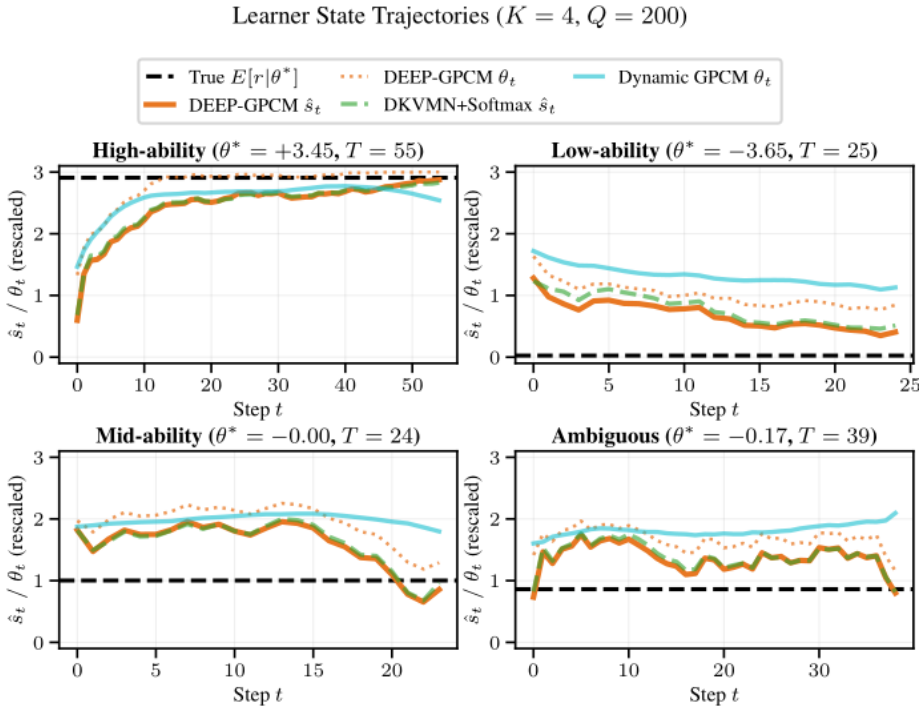

Figure 2: Learner ability trajectories for four archetypes at $K = 4$, $Q = 200$. DEEP-GPCM differentiates high/low ability and growth/decline patterns. DKVMN+Softmax produces identical trajectories, confirming the GPCM head preserves dynamics. Dynamic GPCM fails to differentiate archetypes.

IRT Parameter Recovery and Stability (RQ3)

Neural networks can recover item discrimination without direct supervision. Table 3 shows that DEEP-GPCM achieves strong discrimination recovery ($r_\alpha > 0.70$ for $K \geq 3$), exceeding Static GPCM by factors of 5–7. This represents implicit psychometric inference: the memory network, trained solely to predict ordinal responses, discovers which items are discriminating by observing how responses vary across learners and time. The model never sees ground-truth α during training, yet learns to distinguish informative items from uninformative ones. This finding opens a path toward data-driven psychometrics, where item characteristics emerge as learned representations rather than being imposed through explicit IRT models.

DEEP-GPCM recovers student abilities with correlation $r_\theta = 0.94$ at $K = 4$, approaching the reliability standards of traditional IRT calibration studies. In operational testing, correlations above 0.90 are considered excellent for high-stakes decisions (**?**). This level of recovery enables practical applications: adaptive item selection based on estimated θ_t, diagnostic feedback tied to ability levels, and real-time identification of struggling learners. The gap between DEEP-GPCM ($r_\theta = 0.94$) and Dynamic GPCM

($r_\theta = 0.06$) reveals that distributed memory is essential for ability estimation, not just prediction. Dynamic GPCM's recurrent aggregation captures temporal dependencies but cannot disentangle individual ability from item characteristics. Static GPCM achieves moderate recovery ($r_\theta = 0.58$) by learning population-level ability distributions through fixed embeddings, but cannot track individual learning trajectories over time.

The lower theta recovery for Static and Dynamic GPCM does not reflect deliberate suppression to favor DEEP-GPCM. Rather, these baselines are architecturally optimized for prediction: Static GPCM uses fixed embeddings that prioritize stable item-student associations over dynamic ability tracking, while Dynamic GPCM uses simple recurrent updates that blend ability and item signals. Both achieve strong prediction performance (Table 1) by learning representations that maximize next-response accuracy, but these representations sacrifice parameter identifiability. DEEP-GPCM, by contrast, demonstrates that distributed memory enables both strong prediction and high parameter recovery—the DKVMN architecture naturally separates item characteristics (key memory) from learner state (value memory), providing the representational structure needed for psychometric identifiability without compromising predictive power.

Figures 3 and 4 visualize parameter recovery at $K = 4$, $Q = 200$. Student ability distributions (Figure 3) show that DEEP-GPCM's estimated abilities align closely with ground truth, while Static GPCM captures the overall distribution shape but with reduced precision, and Dynamic GPCM shows minimal correlation. Item parameter scatter plots (Figure 4) confirm strong discrimination recovery ($r_\alpha = 0.707$) and threshold recovery ($\bar{r}_\beta = 0.918$) for DEEP-GPCM.

Table 3: Parameter recovery across models and K at $Q = 200$. DEEP-GPCM recovers discrimination without supervision. Best r_α per K in bold.

	$K = 3$			$K = 4$			$K = 5$			$K = 6$		
Model	r_α	$\bar{r}_\beta$	r_θ	r_α	$\bar{r}_\beta$	r_θ	r_α	$\bar{r}_\beta$	r_θ	r_α	$\bar{r}_\beta$	r_θ
Static GPCM	.700	**.989**	.60	.606	**.983**	.58	.698	**.985**	.62	.539	**.984**	.61
Dynamic GPCM	.645	.976	.05	.700	.966	.06	.694	.965	.06	.523	.960	.07
DEEP-GPCM	**.754**	.816	**.92**	**.707**	.918	**.94**	**.746**	.912	**.95**	**.759**	.927	**.95**

The architectural tradeoff appears in threshold recovery: DEEP-GPCM achieves $\bar{r}_\beta \approx 0.91$ versus 0.98 for Static GPCM. This gap reflects design consequences, not failure. Lookup-based models provide direct gradient paths to β parameters; DEEP-GPCM infers thresholds through memory reads and extraction layers, trading some recovery precision for sequential dynamics modeling. Critically, $r_\beta > 0.91$ across all conditions remains psychometrically meaningful for adaptive testing and item analysis—classical IRT studies routinely accept $r > 0.85$ for operational use (**?**). The key insight is that neural knowledge tracing can serve dual purposes: predicting future performance while recovering interpretable item parameters suitable for high-stakes educational decisions.

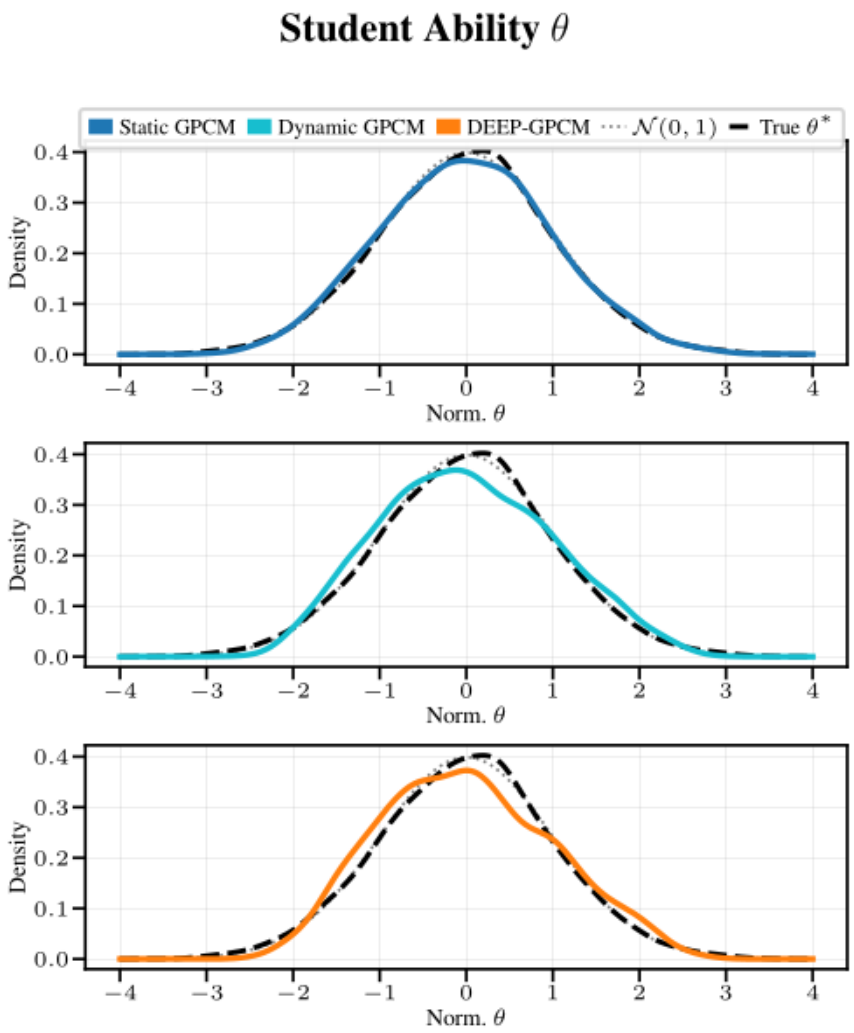

Figure 3: Student ability θ recovery at $K = 4$, $Q = 200$. DEEP-GPCM achieves $r_\theta = 0.94$, approaching operational testing standards. Static GPCM ($r_\theta = 0.58$) captures population distributions but not individual trajectories. Dynamic GPCM ($r_\theta = 0.06$) cannot disentangle ability from item characteristics.

Scalable Item Encoding (RQ4)

Table 4 reveals that encoding strategy and item bank size interact to determine parameter recovery quality. At small scales ($Q = 200$), all three encodings achieve competitive discrimination recovery: SIE leads with $r_\alpha = 0.707$, followed by Separable at 0.739, while LinDecay lags at 0.319. As item banks grow, however, the im-

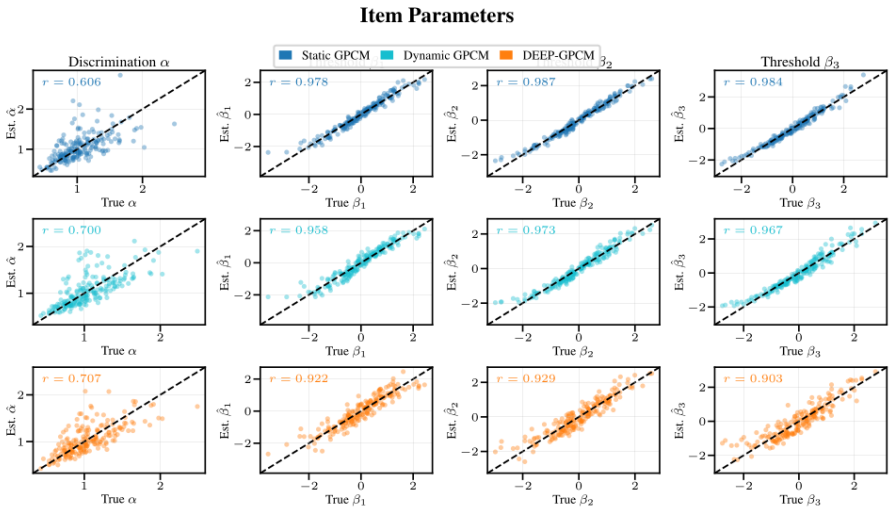

Figure 4: Item parameter recovery (α, β) at $K = 4$, $Q = 200$. DEEP-GPCM achieves $r_\alpha = 0.707$ and $\bar{r}_\beta = 0.918$ through implicit psychometric inference from sequential interaction patterns.

portance of context-aware encoding becomes apparent. At $Q = 2000$, SIE achieves $r_\alpha = 0.856$, demonstrating that frozen random projections maintain item distinctiveness even at large scales. LinDecay shows erratic behavior: poor recovery at $Q = 1000$ ($r_\alpha = 0.198$) but strong recovery at $Q = 2000$ ($r_\alpha = 0.802$), suggesting that the Kronecker product encoding requires sufficient item diversity to learn meaningful discrimination patterns. Separable encoding maintains moderate recovery across scales ($r_\alpha \approx 0.72$–0.80), but SIE consistently achieves the strongest or near-strongest discrimination recovery at every scale.

The key finding is that context-aware encoding becomes increasingly critical as item banks grow. At $Q = 200$, the differences between encodings are modest (0.319–0.739), but at $Q = 2000$, SIE's advantage is clear (0.856 versus 0.773 for Separable). This pattern reflects the architectural properties of frozen random projections: by decoupling item identity from response patterns, SIE prevents gradient crosstalk between items while allowing the response projection to learn shared ordinal structure. The result is 80% parameter reduction compared to LinDecay with maintained or improved recovery quality. For operational testing programs where item banks routinely exceed 1000 questions, SIE enables psychometric validity without architectural bloat.

Theta recovery remains excellent across all encodings and scales ($r_\theta > 0.93$), confirming that ability estimation is robust to encoding choice. Threshold recovery similarly remains high ($\bar{r}_\beta > 0.81$), with SIE achieving the strongest recovery at most scales. Prediction quality (QWK) is stable across encodings, varying by less than 0.02 at any given scale, demonstrating that parameter recovery improvements do not come at the cost of predictive performance.

SIE factorizes item identity and response patterns into independent channels: a

Table 4: Scalability at $K = 4$. SIE achieves strong recovery across all scales. Best r_α per Q in bold.

$Q = 200$					$Q = 500$				
Enc.	**QWK**	r_α	$\bar{r}_\beta$	r_θ	**Enc.**	**QWK**	r_α	$\bar{r}_\beta$	r_θ
LinDecay	.684	.319	.861	**.95**	LinDecay	.697	.419	.818	**.96**
Separable	.684	.739	.872	**.95**	Separable	**.701**	.799	.822	**.95**
SIE	.682	**.707**	**.918**	.94	SIE	.699	**.785**	**.902**	.94
$Q = 1000$					$Q = 2000$				
Enc.	**QWK**	r_α	$\bar{r}_\beta$	r_θ	**Enc.**	**QWK**	r_α	$\bar{r}_\beta$	r_θ
LinDecay	.691	.198	.830	**.95**	LinDecay	.708	.802	.834	**.98**
Separable	**.696**	**.718**	**.913**	.95	Separable	.711	.773	.873	**.98**
SIE	.693	.618	.876	.93	SIE	**.712**	**.856**	**.924**	**.98**

frozen random vector for item identity and a learned projection for ordinal structure. This design eliminates gradient crosstalk between items while allowing the response projection to learn shared structure from all items simultaneously. The result is 80% parameter reduction compared to LinearDecay with maintained recovery quality. For operational testing programs where item banks routinely exceed 1000 questions, efficient encoding enables psychometric validity without architectural bloat.

Ablation and Sensitivity Analysis

Monotonic threshold constraint. Psychometric constraints improve learning. Table 5 shows that enforcing $\beta_0 < \beta_1 < \cdots < \beta_{K-2}$ through softplus gaps improves threshold recovery by 34 points ($\bar{r}_\beta = 0.918$ vs. 0.684). Without this constraint, the model can learn disordered thresholds where some rubric levels are never most probable—a psychometric pathology signaling item dysfunction. This finding challenges the conventional wisdom that constraints reduce model capacity: here, the constraint guides the model toward a more generalizable solution that better captures the data-generating process.

Table 5: Monotonicity constraint ablation at $Q = 200$, $K = 4$. Constraint improves threshold recovery by 34 points. Best per column in bold.

Model	**QWK**	$\boldsymbol{\tau}$	r_α	$\bar{r}_\beta$	r_θ
Monotonic	**.682**	**.513**	**.707**	**.918**	.94
Unconstrained	.679	.510	.676	.684	**.93**

Class Imbalance Robustness (RQ5)

Real educational assessments rarely produce uniformly distributed rubric scores: high-performing students dominate most operational datasets, creating skewed response dis-

Table 6: Class imbalance robustness at $Q = 200$, $K = 4$. Response distribution shows empirical proportions for categories 0–3. QWK and IRT recovery correlations (r_α, r_β, r_θ) are reported on the held-out test set.

Condition	**Distribution**	**QWK**	r_α	r_β	r_θ
Balanced	.31/.21/.20/.29	0.683	0.777	0.927	0.941
Mild	.16/.16/.21/.46	0.669	0.736	0.909	0.931
Severe	.08/.11/.20/.62	0.630	0.641	0.844	0.896
Extreme	.05/.08/.18/.69	0.594	0.543	0.829	0.877

tributions that challenge standard classifiers. We evaluate DEEP-GPCM under three levels of imbalance at $Q = 200$, $K = 4$, induced by shifting the student ability distribution upward ($\theta \sim \mathcal{N}(\mu, 1)$ for $\mu \in \{0.5, 1.0, 1.5\}$), producing mild, severe, and extreme skew toward the highest response category.

Table **??** reveals three findings relevant to practitioners deploying ordinal KT in operational settings.

First, **prediction quality degrades gracefully**. QWK drops from 0.683 under balanced conditions to 0.594 under extreme imbalance—a 13% relative decline even when the dominant category accounts for 69% of responses. The model does not collapse to majority-class prediction, confirming that the WOL's class weighting successfully counteracts distributional skew.

Second, **student ability recovery remains strong**. r_θ decreases from 0.941 to 0.877 across conditions, remaining well above the 0.80 threshold typically considered acceptable for individual-level measurement (**?**). This stability arises because the DKVMN's sequential memory accumulates evidence across interactions regardless of the marginal response distribution, producing reliable per-student ability estimates even when most responses fall in a single category.

Third, **item discrimination is most sensitive to imbalance**. r_α drops from 0.777 to 0.543 under extreme skew. This is expected: when few students respond in the lowest categories, the model has limited information to distinguish how sharply items discriminate at the low end of the ability continuum. Threshold recovery (r_β) is more robust (0.927→0.829) because thresholds are anchored by the ordinal structure of the GPCM even when category frequencies are uneven. For educational applications, this suggests that item parameter calibration should use approximately balanced datasets when possible, while student ability tracking can tolerate substantial distributional skew without meaningful accuracy loss.

Discussion

Educational Implications

DEEP-MIRT extends binary KT to address wider applications for formative assessment practice. By estimating a dynamic ability and categorical probabilities, it differentiates partial knowledge from total misunderstanding, providing richer insights in continu-

ous learning progressions. The model's ability recovery ($r_\theta = 0.94$) meets operational standards for adaptive testing: traditional IRT calibration studies require 20–40 items per student to achieve comparable reliability, but DEEP-GPCM achieves this dynamically from sequential interaction data. This enables real-time ability tracking without requiring students to complete lengthy calibration assessments.

Furthermore, the extracted item parameters describe item-level characteristics beyond immediate predictive use. Unlike classical psychometric calibration, DEEP-MIRT allows online extraction of item characteristics alongside student states, which helps identify underperforming or misaligned items in real time. For example, items with unexpectedly low discrimination ($\alpha < 0.5$) may signal ambiguous wording or construct-irrelevant difficulty, while disordered thresholds indicate rubric levels that are never most probable—a psychometric pathology requiring item revision. This transition from a purely predictive instrument to a measurement-aware system ensures that instructional decisions are grounded in identifiable, theoretically-backed constructs such as discrimination and step difficulty (Heritage, 2007; Wilson, 2005).

Limitations and Future Directions

While the framework establishes a foundation for polytomous knowledge tracing, several limitations present direct avenues for future research.

First, although we proposed DEEP-MIRT as a modular framework with interchangeable components, the current implementation DEEP-GPCM is limited to the GPCM head. A significant research opportunity lies in fully realizing this modularity. For example, a nominal response head using the NRM (Bock, 1972) would extend the framework to unordered polytomous responses such as multiple-choice distractors, enabling misconception-level feedback.

Second, evaluating parameter recovery on synthetic data generated under GPCM assumptions is inherently circular, and proxy signals like attempt counts contain construct-irrelevant variance. Future evaluations will apply the model to authentic polytomous datasets and validate the inferred trajectories against external standardized measures.

Third, while the current backbone effectively maintains a unidimensional student state but struggles with multi-dimensional discrimination recovery. Connecting multi-dimensional instantiations with concept-aligned memory slots would enable per-skill ability tracking, supporting more targeted tracking than a composite θ_t.

Conclusion

The transition from binary to polytomous responses is essential for practical educational assessment. DEEP-MIRT reconciles the predictive performance of sequential neural networks with the rigorness of IRT. By decoupling the sequential backbone from interchangeable psychometric heads, we demonstrate that deep learning can optimize prediction while yielding identifiable measurement constructs. The framework's backward compatibility and scalability confirm its viability as a practical replacement for existing systems. By respecting the ordinal structure of educational data, DEEP-MIRT establishes an alternative for measurement-aware adaptive learning.

References

Baral, S., Botelho, A. F., Erickson, J. A., Benachamardi, P., and Heffernan, N. T. (2021). Improving automated scoring of student open responses in mathematics. In *Proceedings of the 14th International Conference on Educational Data Mining (EDM 2021)*. ERIC ED615565.

Black, P. and Wiliam, D. (1998). Assessment and classroom learning. *Assessment in Education: Principles, Policy & Practice*, 5(1):7–74.

Bock, R. D. (1972). Estimating item parameters and latent ability when responses are scored in two or more nominal categories. *Psychometrika*, 37(1):29–51.

Choi, Y., Lee, Y., Cho, J., Baek, J., Kim, B., Cha, Y., Shin, D., Bae, C., and Heo, J. (2020). Towards an appropriate query, key, and value computation for knowledge tracing. In *Proceedings of the 7th ACM Conference on Learning @ Scale (L@S 2020)*, pages 341–344.

Cohen, J. (1968). Weighted kappa: Nominal scale agreement provision for scaled disagreement or partial credit. *Psychological Bulletin*, 70(4):213–220.

Converse, G., Curi, M., Oliveira, S., and Templin, J. (2021). Incorporating item response theory into knowledge tracing. In *Artificial Intelligence in Education (AIED 2021)*, pages 114–118. Springer.

Corbett, A. T. and Anderson, J. R. (1994). Knowledge tracing: Modeling the acquisition of procedural knowledge. *User Modeling and User-Adapted Interaction*, 4(4):253–278.

Dasgupta, S. and Gupta, A. (2003). An elementary proof of a theorem of Johnson and Lindenstrauss. *Random Structures & Algorithms*, 22(1):60–65.

Embretson, S. E. and Reise, S. P. (2000). *Item Response Theory for Psychologists*. Lawrence Erlbaum Associates, Mahwah, NJ.

Ghosh, A., Heffernan, N., and Lan, A. S. (2020). Context-aware attentive knowledge tracing. arXiv:2007.12324.

Graves, A., Wayne, G., Reynolds, M., Harley, T., Danihelka, I., Grabska-Barwińska, A., Colmenarejo, S. G., Grefenstette, E., Ramalho, T., Agapiou, J., Badia, A. P., Hermann, K. M., Zwols, Y., Ostrovski, G., Cain, A., King, H., Summerfield, C., Blunsom, P., Kavukcuoglu, K., and Hassabis, D. (2016). Hybrid computing using a neural network with dynamic external memory. *Nature*, 538(7626):471–476.

Gutiérrez, P. A., Pérez-Ortiz, M., Sánchez-Monedero, J., Fernández-Navarro, F., and Hervás-Martínez, C. (2016). Ordinal regression methods: Survey and experimental study. *IEEE Transactions on Knowledge and Data Engineering*, 28(1):127–146.

Hambleton, R. and Swaminathan, H. (1985). *Item Response Theory: Principles and Applications*. Springer Netherlands.

Heritage, M. (2007). Formative assessment: What do teachers need to know and do? *Phi Delta Kappan*, 89(2):140–145.

Kane, M. (2013). The argument-based approach to validation. *School Psychology Review*, 42(4):448–457.

Li, X., Luo, F., Ouyang, J., Rojas Pino, L., Li, W., Ding, W., and Gu, C. (2025). A novel deep knowledge tracing model with problem complexity and state stability. *International Journal of Artificial Intelligence in Education*, 35:3270–3293.

Lin, T.-Y., Goyal, P., Girshick, R., He, K., and Dollár, P. (2020). Focal loss for dense object detection. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(2):318–327.

Liu, Z., Liu, Q., Chen, J., Huang, S., and Luo, J. (2023). simpleKT: A simple but tough-to-beat baseline for knowledge tracing. In *Proceedings of the 11th International Conference on Learning Representations (ICLR 2023)*.

Lu, G., Niu, K., Peng, X., Zhou, Y., Zhang, K., and Tai, W. (2024). Self-kt: Self-attentive knowledge tracing with feature fusion pre-training in online education. In *International Joint Conference on Neural Networks (IJCNN)*.

Masters, G. N. and Wright, B. D. (1997). The partial credit model. In van der Linden, W. J. and Hambleton, R. K., editors, *Handbook of Modern Item Response Theory*, pages 101–121. Springer, New York, NY.

Messick, S. (1995). Validity of psychological assessment: Validation of inferences from persons' responses and performances as scientific inquiry into score meaning. *American Psychologist*, 50(9):741–749.

Meulders, M. and Xie, Y. (2004). Person-by-item predictors. In De Boeck, P. and Wilson, M., editors, *Explanatory Item Response Models: A Generalized Linear and Nonlinear Approach*, pages 213–240. Springer, New York, NY.

Mukherjee, K., Khare, A., and Verma, A. (2019). A simple dynamic learning rate tuning algorithm for automated training of dnns. arXiv:1910.11605.

Muraki, E. (1997). A generalized partial credit model. In van der Linden, W. J. and Hambleton, R. K., editors, *Handbook of Modern Item Response Theory*, pages 153–164. Springer, New York, NY.

National Research Council (2001). *Knowing What Students Know: The Science and Design of Educational Assessment*. National Academies Press, Washington, DC.

Pandey, S. and Karypis, G. (2019). A self-attentive model for knowledge tracing. arXiv:1907.06837.

Piech, C., Bassen, J., Huang, J., Ganguli, S., Sahami, M., Guibas, L., and Sohl-Dickstein, J. (2015). Deep knowledge tracing. In *Proceedings of the 29th International Conference on Neural Information Processing Systems (NIPS 2015)*, pages 505–513, Cambridge, MA, USA. MIT Press.

Samejima, F. (1969). Estimation of latent ability using a response pattern of graded scores. *Psychometrika Monograph Supplement*, 34(4, Pt. 2):1–97.

Shute, V. J. (2011). Stealth assessment in computer-based games to support learning. *Computer Games and Instruction*, 55(2):503–524.

Wang, Y. and Heffernan, N. (2013). Extending knowledge tracing to allow partial credit: Using continuous versus binary nodes. In Lane, H. C., Yacef, K., Mostow, J., and Pavlik, P., editors, *Artificial Intelligence in Education (AIED 2013)*, pages 181–188, Berlin, Heidelberg. Springer.

Wilson, M. (2005). *Constructing Measures: An Item Response Modeling Approach*. Lawrence Erlbaum Associates, Mahwah, NJ.

Xu, F., Chen, K., Zhong, M., Liu, L., Liu, H., Luo, X., and Zheng, L. (2024). Dkvmn&mri: A new deep knowledge tracing model based on dkvmn incorporating multi-relational information. *PLOS ONE*, 19(10):1–21.

Yeung, C.-K. (2019). Deep-irt: Make deep learning based knowledge tracing explainable using item response theory. arXiv:1904.11738.

Zhang, J., Shi, X., King, I., and Yeung, D.-Y. (2017). Dynamic key-value memory networks for knowledge tracing. arXiv:1611.08108.